

```
library(tidyverse)

## Warning: package 'tidyverse' was built under R version 4.5.2

## Warning: package 'stringr' was built under R version 4.5.2

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr     1.1.4     v readr     2.1.5
## vforcats   1.0.1     v stringr   1.5.2
## v ggplot2   4.0.0     v tibble    3.3.0
## v lubridate 1.9.4     v tidyverse 1.3.1
## v purrr    1.1.0
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(caret)

## Warning: package 'caret' was built under R version 4.5.2

## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift

library(rpart)
library(rpart.plot)

## Warning: package 'rpart.plot' was built under R version 4.5.2

library(ROCR)

## Warning: package 'ROCR' was built under R version 4.5.2

library(pROC)

## Warning: package 'pROC' was built under R version 4.5.2

## Type 'citation("pROC")' for a citation.
##
## Attaching package: 'pROC'
##
## The following objects are masked from 'package:stats':
##
##     cov, smooth, var
```

```

library(corrplot)

## Warning: package 'corrplot' was built under R version 4.5.2

## corrplot 0.95 loaded

library(janitor)

## Warning: package 'janitor' was built under R version 4.5.2

## 
## Attaching package: 'janitor'
##
## The following objects are masked from 'package:stats':
## 
##     chisq.test, fisher.test

set.seed(123)

getwd()

## [1] "C:/Users/charh/Desktop"

#Load the data

heart <- read.table("processed.cleveland.data",
                     sep = ",",
                     header = FALSE,
                     na.strings = "?")

#Give data column names

colnames(heart) <- c(
  "age", "sex", "cp", "trestbps", "chol",
  "fbs", "restecg", "thalach", "exang",
  "oldpeak", "slope", "ca", "thal", "num"
)

#Convert target variable to binary

heart <- heart %>%
  mutate(
    disease = ifelse(num > 0, 1, 0),
    disease = factor(disease)
  )

#Drop multi-class col

```

```
heart$num <- NULL

glimpse(heart)

## #> Rows: 303
## #> Columns: 14
## #> 
## #> $ age      <dbl> 63, 67, 67, 37, 41, 56, 62, 57, 63, 53, 57, 56, 56, 44, 52, 5-
## #> $ sex      <dbl> 1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1-
## #> $ cp       <dbl> 1, 4, 4, 3, 2, 2, 4, 4, 4, 4, 4, 2, 3, 2, 3, 3, 2, 4, 3, 2, 1-
## #> $ trestbps <dbl> 145, 160, 120, 130, 130, 120, 140, 120, 130, 140, 140, 140, 1-
## #> $ chol      <dbl> 233, 286, 229, 250, 204, 236, 268, 354, 254, 203, 192, 294, 2-
## #> $ fbs       <dbl> 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0-
## #> $ restecg   <dbl> 2, 2, 2, 0, 2, 0, 2, 2, 0, 2, 2, 0, 2, 2, 0, 0, 0, 0, 0, 0, 2-
## #> $ thalach   <dbl> 150, 108, 129, 187, 172, 178, 160, 163, 147, 155, 148, 153, 1-
## #> $ exang     <dbl> 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1-
## #> $ oldpeak   <dbl> 2.3, 1.5, 2.6, 3.5, 1.4, 0.8, 3.6, 0.6, 1.4, 3.1, 0.4, 1.3, 0-
## #> $ slope     <dbl> 3, 2, 2, 3, 1, 1, 3, 1, 2, 3, 2, 2, 2, 1, 1, 1, 3, 1, 1, 1, 2-
## #> $ ca        <dbl> 0, 3, 2, 0, 0, 0, 2, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0-
## #> $ thal      <dbl> 6, 3, 7, 3, 3, 3, 3, 3, 7, 7, 6, 3, 6, 7, 7, 3, 7, 3, 3, 3, 3-
## #> $ disease   <fct> 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0-
```

```
colSums(is.na(heart))
```

```

##      age      sex      cp trestbps      chol      fbs restecg thalach
##      0        0        0        0        0        0        0        0
##      exang    oldpeak slope      ca      thal disease
##      0        0        0        4        2        0

```

#Median imputation for numeric #Mode for categorical

```
for (col in names(heart)) {  
  if (is.numeric(heart[[col]])) {  
    heart[[col]][is.na(heart[[col]])] <- median(heart[[col]], na.rm = TRUE)  
  } else {  
    heart[[col]][is.na(heart[[col]])] <- names(which.max(table(heart[[col]])))  
  }  
}
```

#Summary Stats

`summary(heart)`

```

##      age          sex          cp          trestbps
##  Min.   :29.00   Min.   :0.0000   Min.   :1.000   Min.   : 94.0
##  1st Qu.:48.00   1st Qu.:0.0000   1st Qu.:3.000   1st Qu.:120.0
##  Median :56.00   Median :1.0000   Median :3.000   Median :130.0
##  Mean    :54.44   Mean    :0.6799   Mean    :3.158   Mean    :131.7
##  3rd Qu.:61.00   3rd Qu.:1.0000   3rd Qu.:4.000   3rd Qu.:140.0
##  Max.   :77.00   Max.   :1.0000   Max.   :4.000   Max.   :200.0
##      chol          fbs          restecg          thalach
##  Min.   :126.0   Min.   :0.0000   Min.   :0.0000   Min.   : 71.0

```

```

## 1st Qu.:211.0 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:133.5
## Median :241.0 Median :0.0000 Median :1.0000 Median :153.0
## Mean   :246.7 Mean   :0.1485 Mean   :0.9901 Mean   :149.6
## 3rd Qu.:275.0 3rd Qu.:0.0000 3rd Qu.:2.0000 3rd Qu.:166.0
## Max.   :564.0 Max.   :1.0000 Max.   :2.0000 Max.   :202.0
##      exang      oldpeak      slope      ca
## Min.   :0.0000  Min.   :0.00  Min.   :1.000  Min.   :0.0000
## 1st Qu.:0.0000 1st Qu.:0.00 1st Qu.:1.000 1st Qu.:0.0000
## Median :0.0000  Median :0.80  Median :2.000  Median :0.0000
## Mean   :0.3267  Mean   :1.04  Mean   :1.601  Mean   :0.6634
## 3rd Qu.:1.0000 3rd Qu.:1.60 3rd Qu.:2.000 3rd Qu.:1.0000
## Max.   :1.0000  Max.   :6.20  Max.   :3.000  Max.   :3.0000
##      thal      disease
## Min.   :3.000 0:164
## 1st Qu.:3.000 1:139
## Median :3.000
## Mean   :4.723
## 3rd Qu.:7.000
## Max.   :7.000

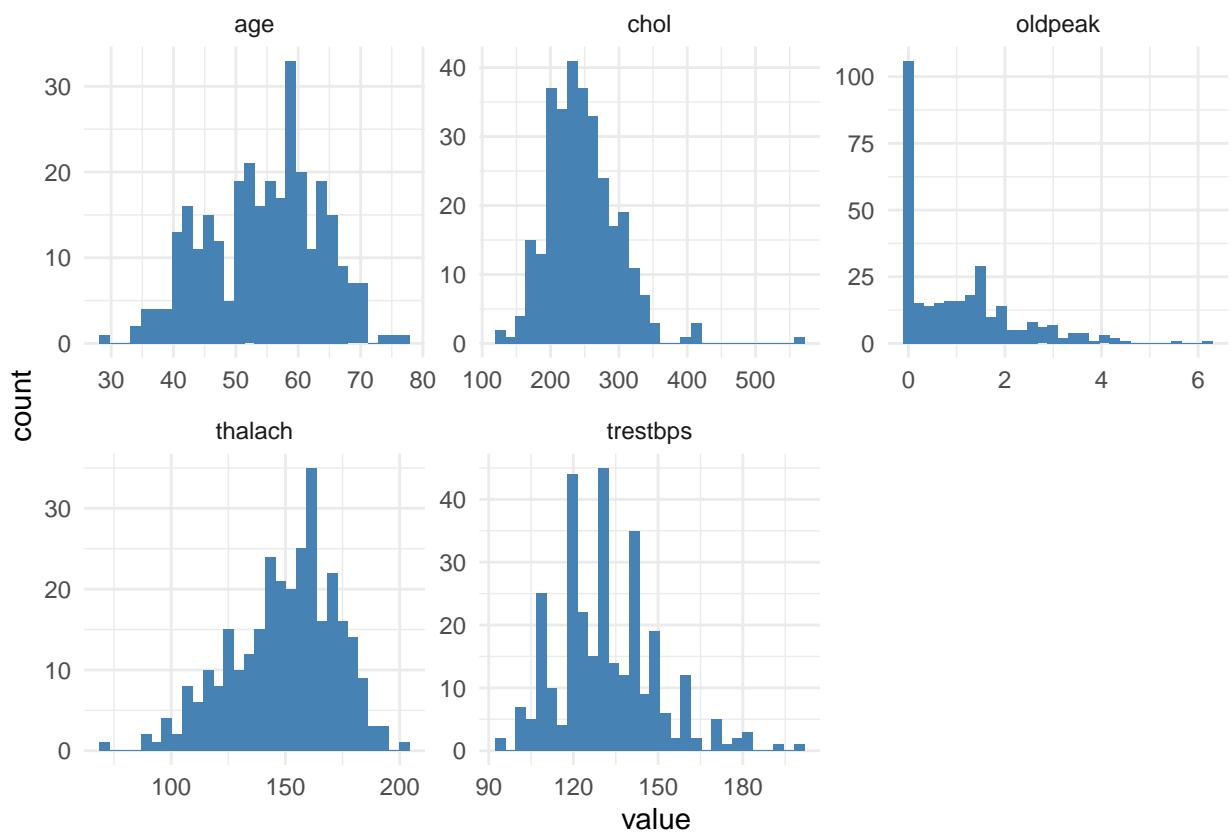
```

#Histogram of numeric columns

```

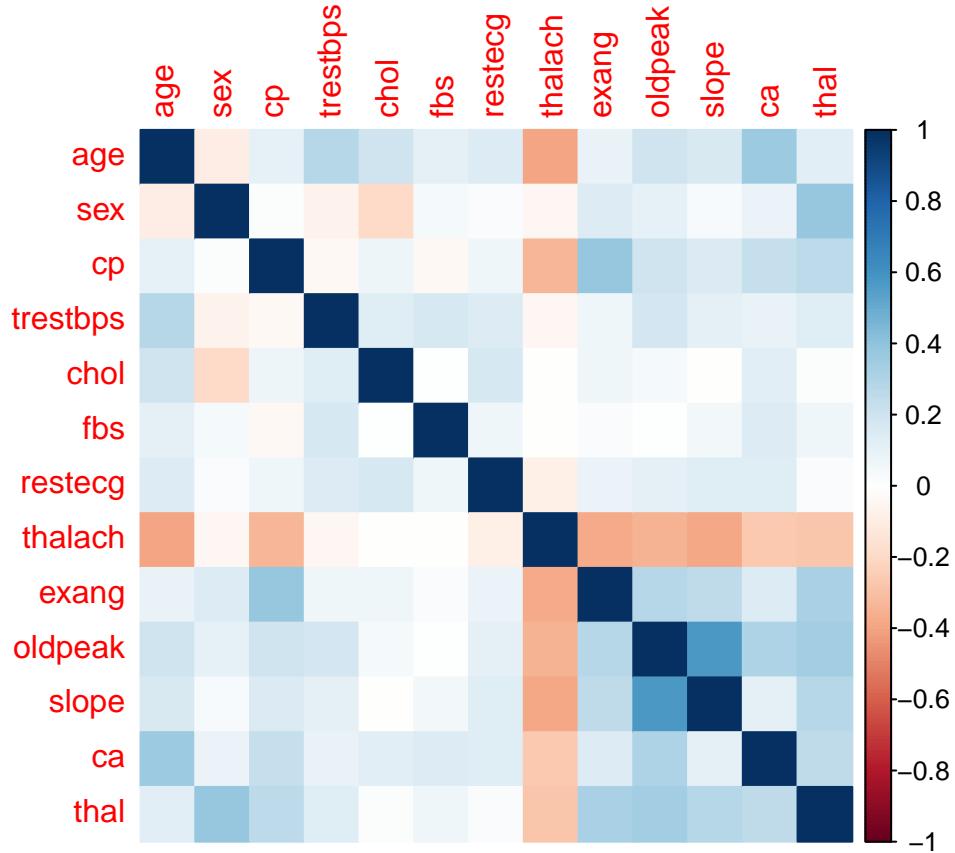
heart %>%
  select(age, trestbps, chol, thalach, oldpeak) %>%
  gather() %>%
  ggplot(aes(value)) +
  geom_histogram(bins = 30, fill = "steelblue") +
  facet_wrap(~ key, scales = "free") +
  theme_minimal()

```



#Correlation plot

```
numeric_vars <- heart %>% select_if(is.numeric)
corrplot(cor(numeric_vars), method = "color")
```

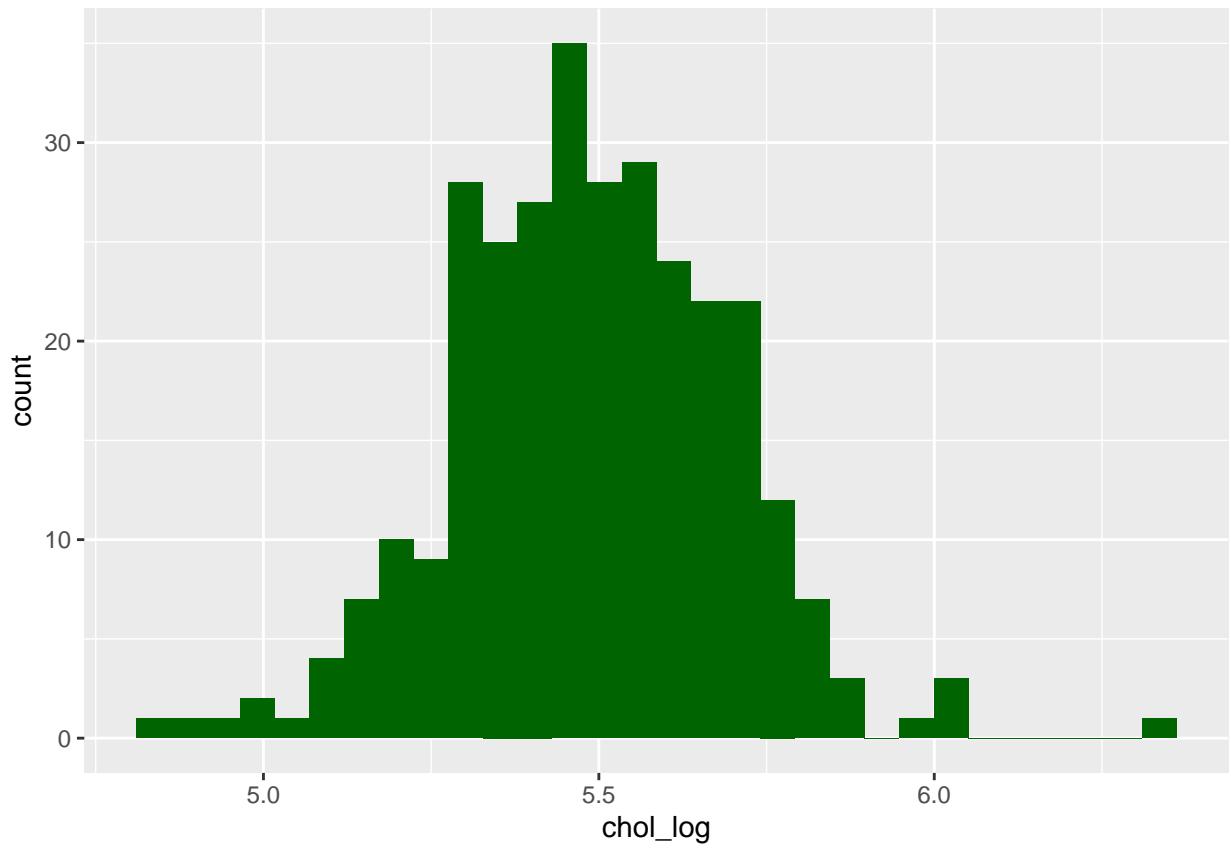


```

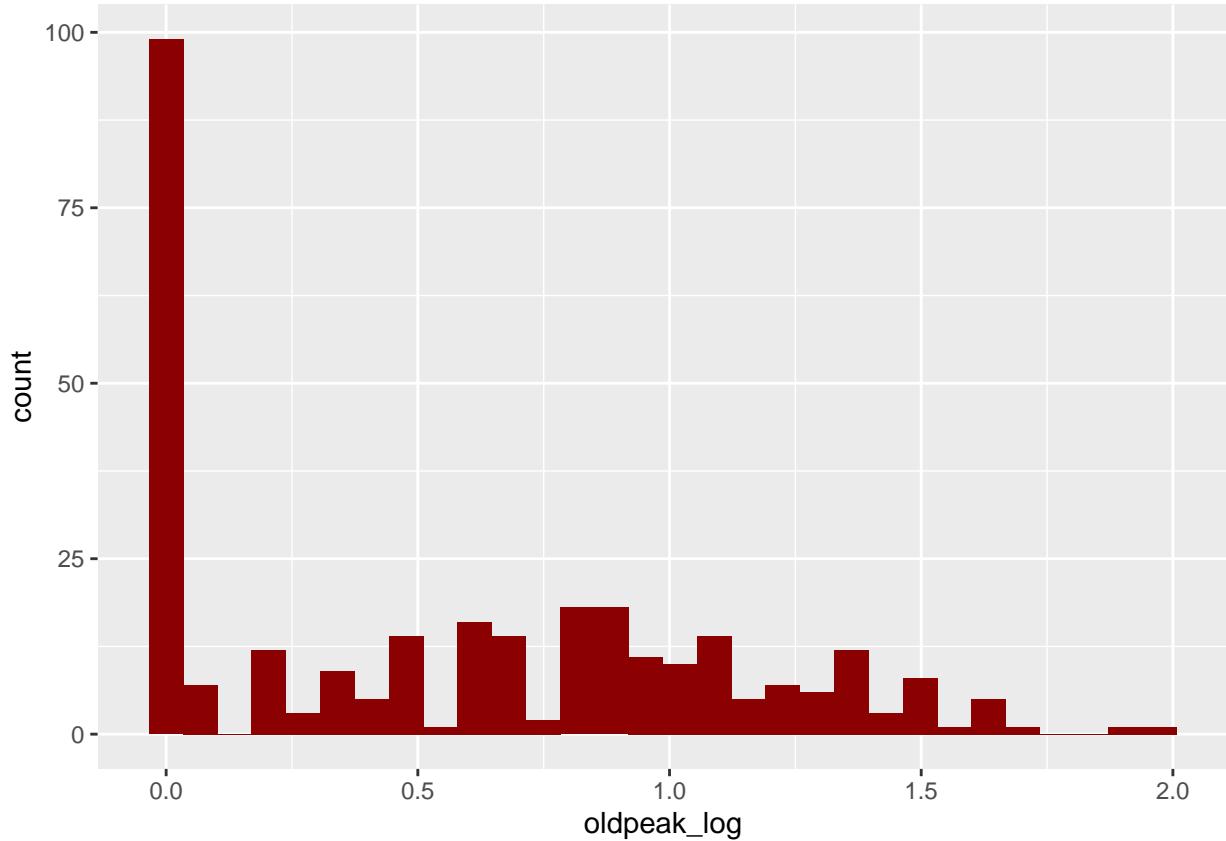
heart <- heart %>%
  mutate(
    chol_log = log(chol),
    oldpeak_log = log(oldpeak + 1) # avoid log(0)
  )

#Check transformed distributions
ggplot(heart, aes(chol_log)) + geom_histogram(bins = 30, fill = "darkgreen")

```



```
ggplot(heart, aes(oldpeak_log)) + geom_histogram(bins = 30, fill = "darkred")
```



```
train_index <- createDataPartition(heart$disease, p = 0.8, list = FALSE)
train <- heart[train_index, ]
test <- heart[-train_index, ]
```

```
log_model <- glm(disease ~ age + sex + cp + trestbps + chol_log +
                    thalach + exang + oldpeak_log + slope + ca + thal,
                    data = train,
                    family = binomial)
```

```
summary(log_model)
```

```
##
## Call:
## glm(formula = disease ~ age + sex + cp + trestbps + chol_log +
##       thalach + exang + oldpeak_log + slope + ca + thal, family = binomial,
##       data = train)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -10.720695   6.223980 -1.722 0.084982 .
## age         -0.009907   0.026072 -0.380 0.703957
## sex          1.617671   0.535150  3.023 0.002504 **
## cp           0.596877   0.211163  2.827 0.004704 **
## trestbps     0.021673   0.011497  1.885 0.059419 .
## chol_log     1.037735   1.063548  0.976 0.329199
```

```

## thalach      -0.025793  0.011728 -2.199 0.027862 *
## exang        1.078295  0.443475  2.431 0.015038 *
## oldpeak_log  0.536189  0.487583  1.100 0.271469
## slope        0.489037  0.373771  1.308 0.190742
## ca           0.985229  0.257613  3.824 0.000131 ***
## thal         0.279325  0.105258  2.654 0.007961 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 336.61  on 243  degrees of freedom
## Residual deviance: 173.15  on 232  degrees of freedom
## AIC: 197.15
##
## Number of Fisher Scoring iterations: 6

exp(cbind(OR = coef(log_model), confint(log_model)))

## Waiting for profiling to be done...

##          OR      2.5 %    97.5 %
## (Intercept) 2.208316e-05 8.778553e-11 4.2033738
## age         9.901421e-01 9.402760e-01 1.0420613
## sex         5.041335e+00 1.827108e+00 15.1094670
## cp          1.816438e+00 1.212642e+00 2.7931559
## trestbps   1.021910e+00 9.993478e-01 1.0457683
## chol_log    2.822817e+00 3.521370e-01 23.7091671
## thalach    9.745369e-01 9.515017e-01 0.9966345
## exang      2.939662e+00 1.234067e+00 7.1005581
## oldpeak_log 1.709480e+00 6.582892e-01 4.5028508
## slope      1.630746e+00 7.760973e-01 3.3982154
## ca          2.678426e+00 1.657918e+00 4.5810995
## thal       1.322237e+00 1.076489e+00 1.6298324

#Predictions on test set
log_prob <- predict(log_model, test, type = "response")
log_pred <- ifelse(log_prob > 0.5, 1, 0)

#Confusion matrix
confusionMatrix(factor(log_pred), test$disease)

## Confusion Matrix and Statistics
##
##          Reference
## Prediction 0 1
##          0 28 5
##          1 4 22
##
##          Accuracy : 0.8475
##          95% CI : (0.7301, 0.9278)
## No Information Rate : 0.5424

```

```

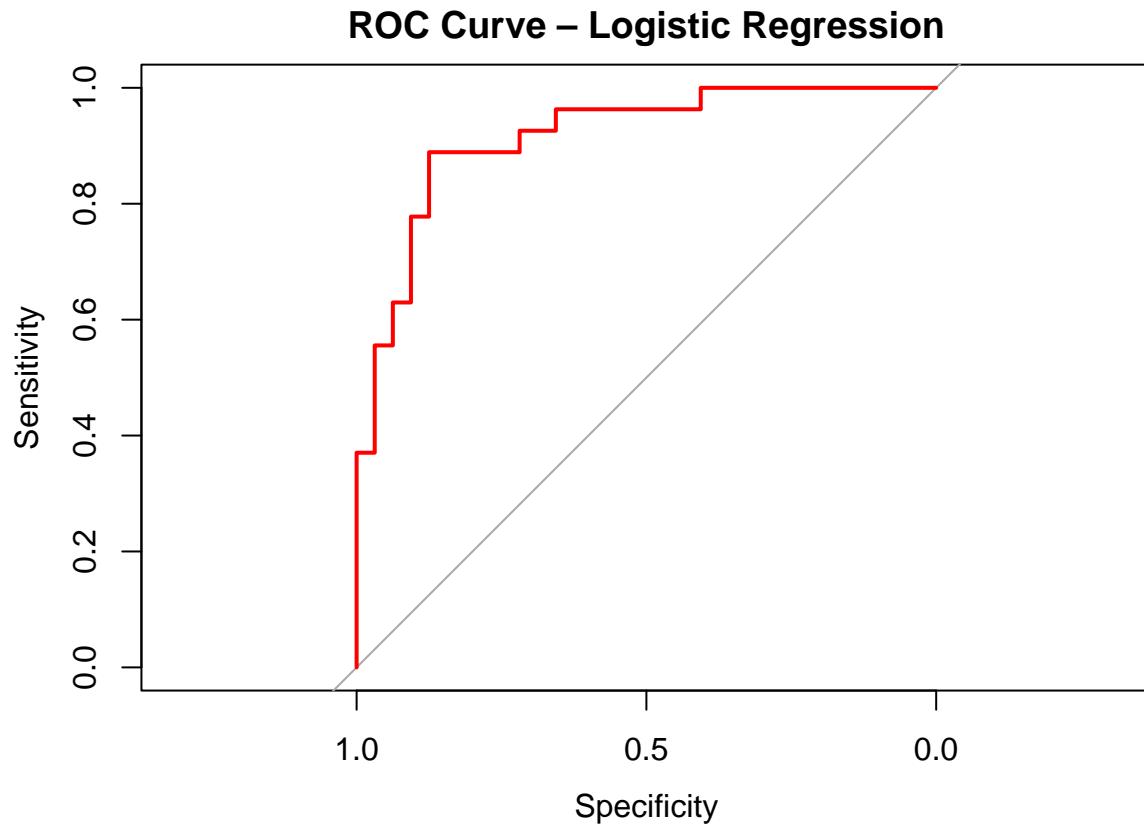
##      P-Value [Acc > NIR] : 7.195e-07
##
##          Kappa : 0.6918
##
##  McNemar's Test P-Value : 1
##
##          Sensitivity : 0.8750
##          Specificity : 0.8148
##          Pos Pred Value : 0.8485
##          Neg Pred Value : 0.8462
##          Prevalence : 0.5424
##          Detection Rate : 0.4746
##          Detection Prevalence : 0.5593
##          Balanced Accuracy : 0.8449
##
##          'Positive' Class : 0
##

#ROC curve & AUC
roc_obj <- roc(test$disease, log_prob)

## Setting levels: control = 0, case = 1
## Setting direction: controls < cases

plot(roc_obj, col = "red", main = "ROC Curve - Logistic Regression")

```



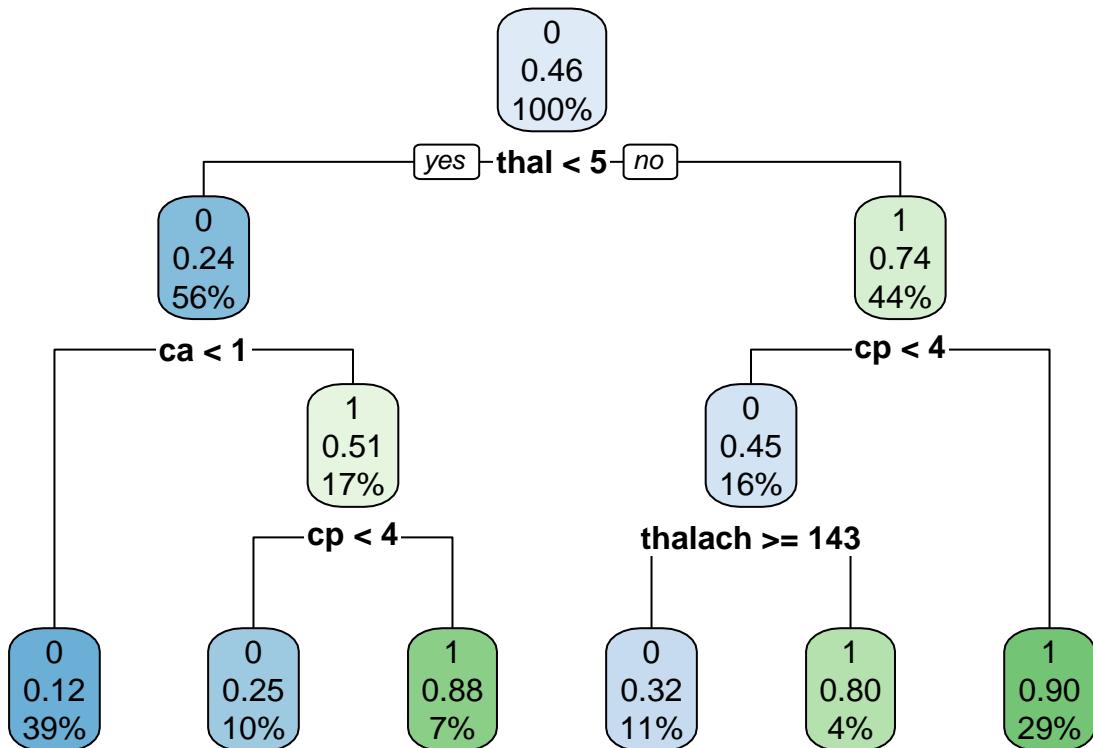
```
auc(roc_obj)
```

```
## Area under the curve: 0.9167
```

```
tree_model <- rpart(disease ~ age + sex + cp + trestbps + chol_log +
                      thalach + exang + oldpeak_log + slope + ca + thal,
                      data = train,
                      method = "class",
                      cp = 0.01)
```

```
#Visualize tree
```

```
rpart.plot(tree_model, type = 2, extra = 106)
```



```
#Predictions & evaluation
```

```
tree_pred <- predict(tree_model, test, type = "class")
confusionMatrix(tree_pred, test$disease)
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction 0 1
##          0 28 6
##          1  4 21
##          
```

```

##          Accuracy : 0.8305
## 95% CI : (0.7103, 0.9156)
## No Information Rate : 0.5424
## P-Value [Acc > NIR] : 3.14e-06
##
##          Kappa : 0.6566
##
## McNemar's Test P-Value : 0.7518
##
##          Sensitivity : 0.8750
##          Specificity : 0.7778
## Pos Pred Value : 0.8235
## Neg Pred Value : 0.8400
##          Prevalence : 0.5424
## Detection Rate : 0.4746
## Detection Prevalence : 0.5763
## Balanced Accuracy : 0.8264
##
## 'Positive' Class : 0
##

```

```

#ROC & AUC for tree
tree_prob <- predict(tree_model, test)[, 2]
roc_tree <- roc(test$disease, tree_prob)

```

```

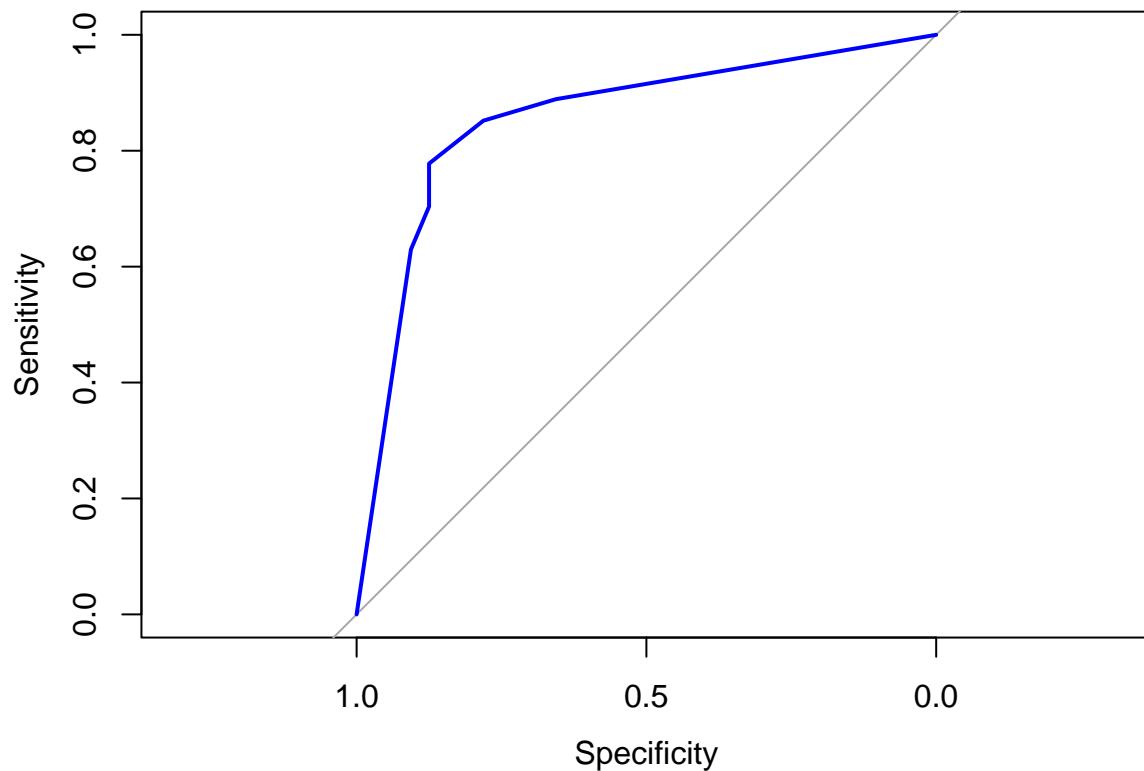
## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

plot(roc_tree, col = "blue", main = "ROC Curve - Decision Tree")

```

ROC Curve – Decision Tree



```
auc(roc_tree)

## Area under the curve: 0.8553

train_control <- trainControl(
  method = "cv",
  number = 5,
  classProbs = TRUE,
  summaryFunction = twoClassSummary
)

heart$disease <- factor(heart$disease, levels = c(0, 1), labels = c("No", "Yes"))

# Check levels
levels(heart$disease)

## [1] "No"   "Yes"

cv_log_model <- train(
  disease ~ age + sex + cp + trestbps + chol_log +
    thalach + exang + oldpeak_log + slope + ca + thal,
  data = heart,
  method = "glm",
  family = binomial,
```

```

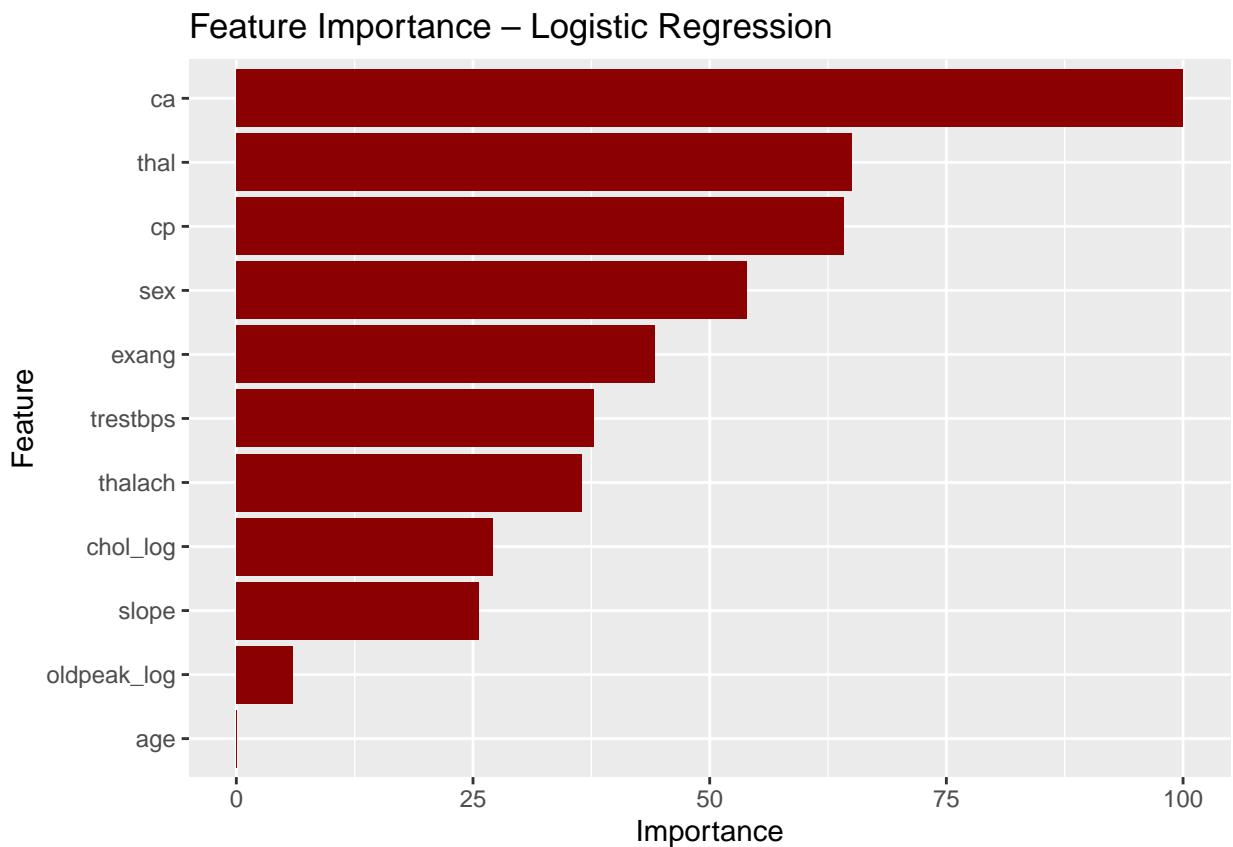
    metric = "ROC",
    trControl = train_control
)

cv_log_model

## Generalized Linear Model
##
## 303 samples
## 11 predictor
## 2 classes: 'No', 'Yes'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 242, 242, 243, 243, 242
## Resampling results:
##
## ROC      Sens      Spec
## 0.8996139 0.8784091 0.805291

varImp(cv_log_model) %>%
  ggplot(aes(x = Overall, y = reorder(rownames(.), Overall))) +
  geom_col(fill = "darkred") +
  labs(title = "Feature Importance - Logistic Regression")

```



```

results <- data.frame(
  Model = c("Logistic Regression", "Decision Tree"),
  Accuracy = c(
    confusionMatrix(factor(log_pred), factor(test$disease))$overall["Accuracy"],
    confusionMatrix(tree_pred, test$disease)$overall["Accuracy"]
  ),
  AUC = c(
    auc(roc_obj),
    auc(roc_tree)
  )
)

results

##           Model  Accuracy      AUC
## 1 Logistic Regression 0.8474576 0.9166667
## 2      Decision Tree 0.8305085 0.8553241

```